

Introduction to the bartMachine R package

Saint Louis R User Group

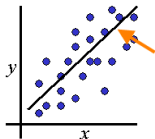
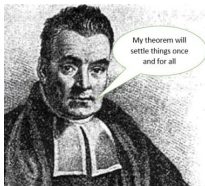
John Snyder

April 11, 2019

What is BART?

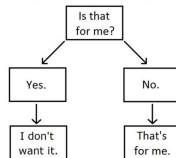


Bayesian Additive Regression Trees



regression
line

My Cat's Decision-Making Tree.



How it works

- ▶ Ensemble method which is the sum of many shallow trees.
- ▶ Complexity is regularized via Bayesian “priors.”
 - ▶ This frees us from ad hoc decisions
- ▶ Uses “Bayesian Backfitting”
 - ▶ Each tree is sequentially exposed to the residuals when all other trees are used to predict

Results:

- ▶ Each tree describes a tiny amount of the structure
- ▶ The Bayesian structure means variation is fully quantified.
 - ▶ Intervals, p-values, and model selection become possible
- ▶ Outperforms many common models in out of sample prediction.

Powerful Predictive Performance

- ▶ Test RMSE of 100 random datasets simulated from various nonlinear functions (added noise with $s=1$)

Function	BART	XGBoost*	Random Forest*	Linear Reg(lol)
Friedman	1.08	1.21	1.64	2.61
Mirsha's Bird	1.53	2.78	2.90	26.59
Weird Exp	1.04	1.05	1.07	6.08
Linear	1.025	1.032	1.034	1.004

Details in Simulation.R

- ▶ `bartMachine` is relatively unknown
 - ▶ `xgboost`: ~43k downloads per month
 - ▶ `randomForest`: ~88k downloads per month
 - ▶ `bartMachine`: ~2k downloads per month
- ▶ No full featured python library

Package Overview:

- ▶ Based in Java
- ▶ Functions for Cross Validation
- ▶ Model fitting:
 - ▶ Is done in parallel¹
 - ▶ Can incorporate missing data
- ▶ Lots of fun statistical things
 - ▶ Credible interval calculation
 - ▶ Diagnostic plots/tests
- ▶ Model free variable selection
- ▶ Interaction detection
- ▶ Export fit trees

¹MCMC sampling is used, so speedups during model fitting aren't great

Installation and loading steps

1. Google “How to install rJava on [your OS]”
2. Do that
3. Run the following

```
install.packages("bartMachine")
```

To load the package with:

- ▶ 10GB of memory
- ▶ All but one core available for compute

```
options(java.parameters = "-Xmx10g")  
library(bartMachine)  
numcores <- parallel::detectCores()  
set_bart_machine_num_cores(numcores - 1)
```

Boston Data

We will consider the Boston dataset from the MASS package

```
data(Boston)
Boston %>% round(digits=2) %>% head
```

```
##      crim zn  indus  chas  nox   rm  age  dis rad tax ptratio  black lstat medv
## 1 0.01 18   2.31    0 0.54 6.58 65.2 4.09   1 296    15.3 396.90  4.98 24.0
## 2 0.03  0   7.07    0 0.47 6.42 78.9 4.97   2 242    17.8 396.90  9.14 21.6
## 3 0.03  0   7.07    0 0.47 7.18 61.1 4.97   2 242    17.8 392.83  4.03 34.7
## 4 0.03  0   2.18    0 0.46 7.00 45.8 6.06   3 222    18.7 394.63  2.94 33.4
## 5 0.07  0   2.18    0 0.46 7.15 54.2 6.06   3 222    18.7 396.90  5.33 36.2
## 6 0.03  0   2.18    0 0.46 6.43 58.7 6.06   3 222    18.7 394.12  5.21 28.7
```

- ▶ Target: Predict median home value(medv)
- ▶ Features include:
 1. lstat: Proportion of low income individuals
 2. rm: Average number of rooms per dwelling
 3. age: proportion of old homes
 4. etc. . .

Fitting BART

```
y <- Boston$medv
X <- Boston %>% dplyr::select(-"medv")

#Fit BART model
bart.model <- bartMachine(X,y,
                           num_trees = 200,
                           num_burn_in = 1000,
                           num_iterations_after_burn_in = 5000)
```

BART is fit with MCMC, which requires a “burnin” set of initial iterations which are discarded.

- ▶ NOTE: with 10 cores, each thread would fit $1000 + 5000/10$ or 1500 iterations

The BART object

```
> bart.model
bartMachine v1.2.3 for regression

training data n = 506 and p = 13
built in 19.4 secs on 31 cores, 200 trees, 1000 burn-in and 5000 post.

sigsq est for y beforehand: 21.938
avg sigsq estimate after burn-in: 2.70319

in-sample statistics:
  L1 = 430.35
  L2 = 654.72
  rmse = 1.14
  Pseudo-Rsq = 0.9847
p-val for shapiro-wilk test of normality of residuals: 0
p-val for zero-mean noise: 0.98467
```

Did we overfit?

K fold CV using our fit model

```
k_fold_cv(X, y, k_folds = 10,  
          num_trees = 200,  
          num_burn_in = 1000,  
          num_iterations_after_burn_in = 5000)
```

```
.....
```

```
$y_hat
```

```
[1] 24.932807 20.522438 31.149051 36.071931...
```

```
$L1_err
```

```
[1] 1022.238
```

```
$L2_err
```

```
[1] 4624.004
```

```
$rmse
```

```
[1] 3.02297
```

```
$PseudoRsq
```

```
[1] 0.8917508
```

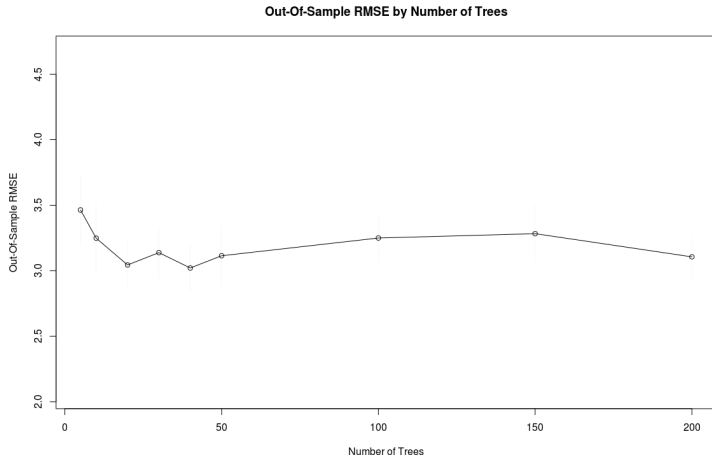
```
$folds
```

```
[1] 8 10 3 1 ....
```

Select the number of Trees

```
rmse_by_num_trees(bart.model, num_replicates = 20)
```

- ▶ This fits the model 20 times for a number of various numbers of trees
- ▶ Aggregates out of sample RMSE



Cross Validate

```
bart.model.cv <- bartMachineCV(X, y,  
                                num_burn_in = 1000,  
                                num_iterations_after_burn_in = 5000)
```

- ▶ To automatically cross validate, we just need to add a CV to the end of the function!
- ▶ Will fit a variety of models and return the best one.
 - ▶ Similar to the caret package
- ▶ This takes a while.
- ▶ `bart.model.cv` will be the model object for the best fit

Cross Validate

```
> bart.model.cv$cv_stats
```

	k	nu	q	num_trees	oos_error	% diff with lowest
[1,]	2	3	0.90	50	2.860489	0.000000
[2,]	3	3	0.90	50	2.922613	2.171807
[3,]	2	3	0.99	200	2.924967	2.254080
[4,]	2	3	0.90	200	2.933799	2.562845
[5,]	2	10	0.75	200	2.956046	3.340573
[6,]	2	10	0.75	50	3.036727	6.161102
[7,]	3	10	0.75	50	3.040292	6.285746
[8,]	3	3	0.99	50	3.060013	6.975171
[9,]	3	3	0.90	200	3.086522	7.901900
[10,]	3	3	0.99	200	3.093631	8.150404
[11,]	3	10	0.75	200	3.123639	9.199475
[12,]	5	3	0.90	50	3.140318	9.782571
[13,]	5	3	0.99	50	3.158856	10.430626
[14,]	5	10	0.75	50	3.215616	12.414893
[15,]	2	3	0.99	50	3.342410	16.847512
[16,]	5	3	0.90	200	3.446951	20.502168
[17,]	5	3	0.99	200	3.447560	20.523454
[18,]	5	10	0.75	200	3.459553	20.942727

Cross Validate

```
k_fold_cv(X, y, k_folds = 5,  
          k=2, nu=3, q=.90, num_trees = 50,  
          num_burn_in = 1000,  
          num_iterations_after_burn_in = 5000)
```

.....

\$rmse

[1] 2.92

\$PseudoRsq

[1] 0.9138275

This is better than before!

Variable Importance

```
investigate_var_importance(bart.model.cv)
```

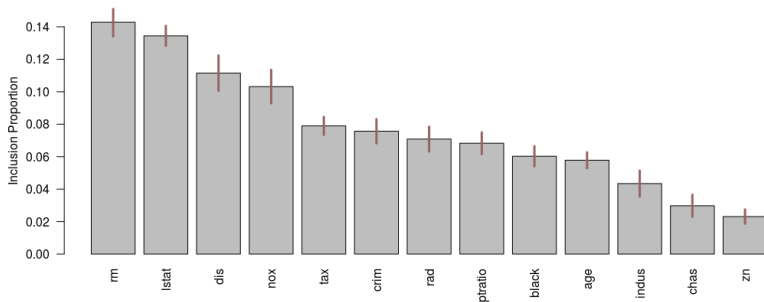


Figure 2

Joint Variable Importance

```
interaction_investigator(bart.model.cv)
```

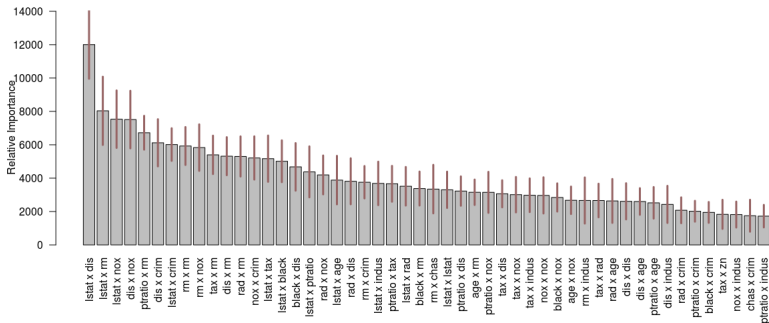


Figure 3

Model Selection

```
VarSel <- var_selection_by_permute(bart.model.cv, bottom_margin = 5)  
VarSel$important_vars_local_names
```

```
[1] "rm"      "lstat"    "dis"      "nox"
```

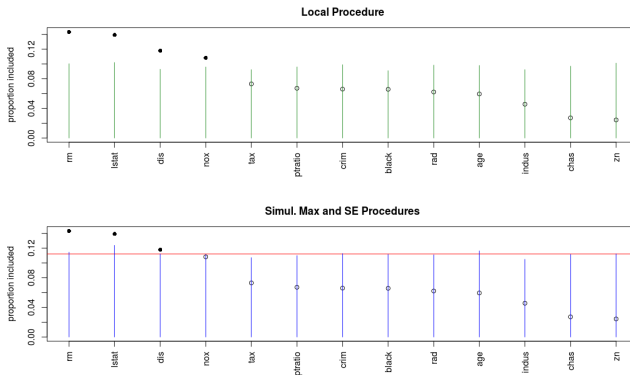


Figure 4

Partial Dependence plots!

```
cov_importance_test(bart.model.cv, covariates = "lstat")  
pd_plot(bart.model.cv, j = "lstat")
```

P-Value = 0

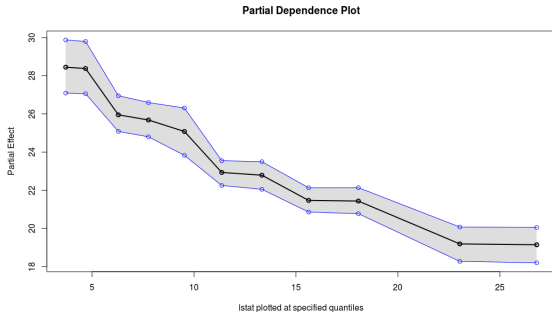


Figure 5

- Low income areas have a negative relationship with home value

Partial Dependence plots!

```
cov_importance_test(bart.model.cv, covariates = "rm")  
pd_plot(bart.model.cv, j = "rm")
```

P-Value = 0

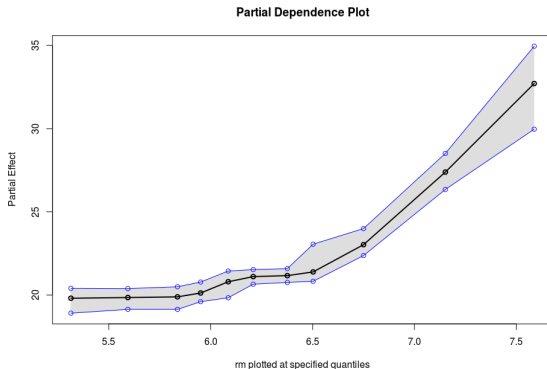


Figure 6

- Number of rooms has a positive relationship with home value


Conclusions supported by regressions

```
RegressMod <- lm(formula = medv ~ ., data = Boston)
summary(RegressMod)
```

```
##
## Call:
## lm(formula = medv ~ ., data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.595  -2.730  -0.518   1.777   26.199
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.646e+01  5.103e+00   7.144 3.28e-12 ***
## crim        -1.080e-01  3.286e-02  -3.287 0.001087 **
## zn           4.642e-02  1.373e-02   3.382 0.000778 ***
## indus        2.056e-02  6.150e-02   0.334 0.738288
## chas         2.687e+00  8.616e-01   3.118 0.001925 **
## nox         -1.777e+01  3.820e+00  -4.651 4.25e-06 ***
## rm           3.810e+00  4.179e-01   9.116 < 2e-16 ***
## age          6.922e-04  1.321e-02   0.052 0.958229
## dis         -1.476e+00  1.995e-01  -7.398 6.01e-13 ***
## rad          3.060e-01  6.635e-02   4.613 5.07e-06 ***
## tax         -1.233e-02  3.760e-03  -3.280 0.001112 **
## ptratio     -9.527e-01  1.308e-01  -7.283 1.31e-12 ***
## black        9.312e-03  2.686e-03   3.467 0.000573 ***
## lstat       -5.248e-01  5.072e-02 -10.347 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.745 on 492 degrees of freedom
## Multiple R-squared:  0.7406, Adjusted R-squared:  0.7338
## F-statistic: 108.1 on 13 and 492 DF, p-value: < 2.2e-16
```

Aside: Partial Dependence

```
nsim <- 1000  
x1 <- runif(nsim,-3.14,3.14)  
x2 <- runif(nsim,-3.14,3.14)  
  
y <- x1 + sin(x2) + rnorm(nsim)  
plot(x2,y)
```



Aside: Partial Dependence

```
bart.model <- bartMachine(data.frame(x1,x2),y,  
                               num_trees = 200,  
                               num_burn_in = 1000,  
                               num_iterations_after_burn_in = 4000)  
pd_plot(bart.model, j = "x2")
```

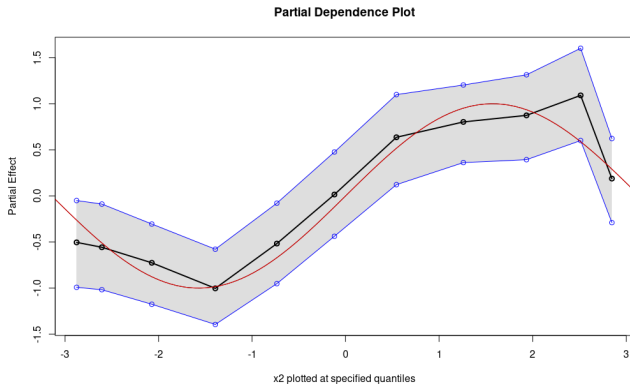


Figure 7

Notable Arguments

The `bartMachine` function has a number of arguments. Default values are listed.

1. `use_missing_data = FALSE`
 - ▶ Uses rows with missing data *without* imputing
2. `mem_cache_for_speed = TRUE`
3. `serialize = FALSE`
 - ▶ Used to store the java object
 - ▶ `save(bart.model, file="MyBart.Rdata")`
4. `cov_prior_vec = NULL`
 - ▶ Place informative priors on variable inclusion
5. If `y` is a factor, classification is performed automatically.

Other Statistical Things: Credible/Prediction intervals!

```
MeanDF <- colMeans(X)
MeanDF %>% head
```

```
      crim      zn  indus  chas   nox    rm   age   dis
1 3.614 11.364 11.137 0.069 0.555 6.285 68.575 3.795
      rad      tax ptratio  black  lstat
1 9.549 408.237 18.456 356.674 12.653
```

```
predict(bart.model.cv, newdata = MeanDF)
calc_credible_intervals(bart.model.cv, MeanDF, ci_conf = 0.95) %>% round(digits=2)
calc_prediction_intervals(bart.model.cv, MeanDF, pi_conf = 0.95) %>% round(digits=2)
```

```
[1] 23.60462
```

```
      ci_lower_bd ci_upper_bd
[1,]          17.69          29.3
```


Other Statistical Things: Assumption Checking

```
check_bart_error_assumptions(bart.model.cv)
```

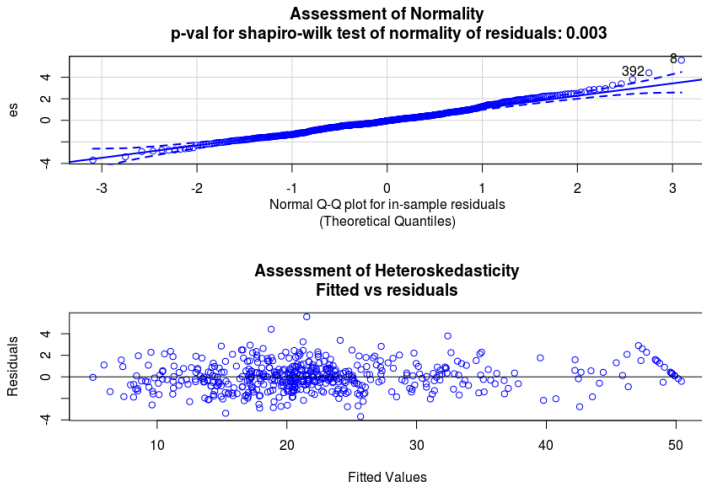
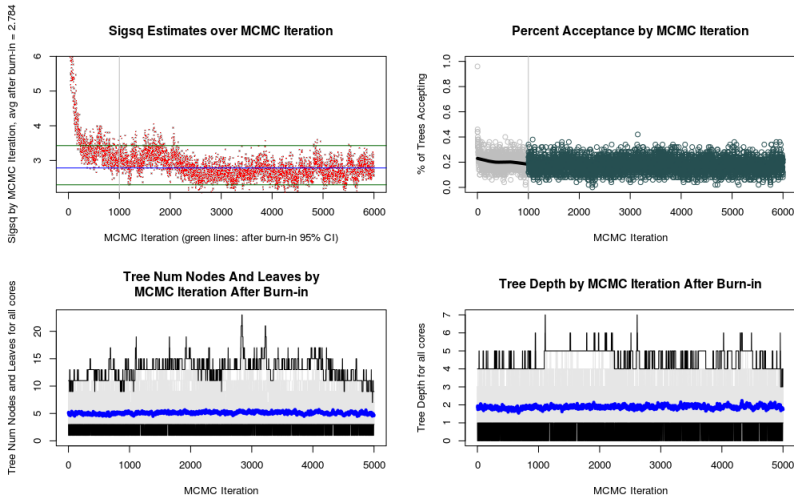


Figure 8

Other Statistical Things: Convergence Diagnostics

```
plot_convergence_diagnostics(bart.model.cv)
```

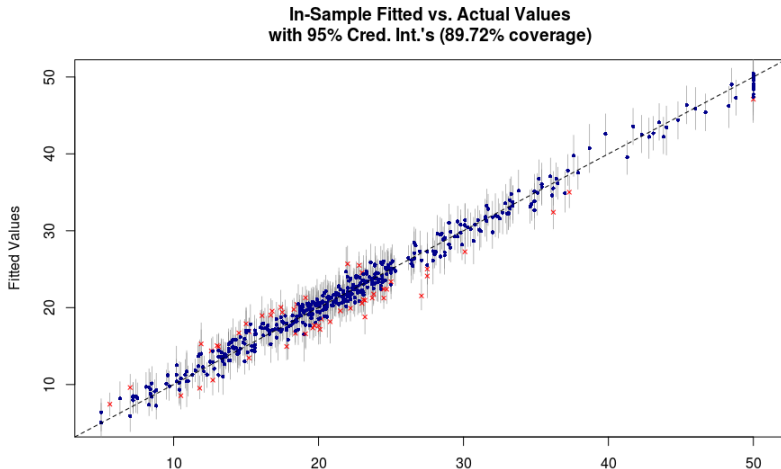
MCMC must “converge”



Other Statistical Things: Observations vs Predictions

```
plot_y_vs_yhat(bart.model.cv, credible_intervals = TRUE)
```

Looking for a 1:1 relationship



John's Final Thought

- ▶ BART is a powerful technique which brings many advantages
 - ▶ At the expense of computational efficiency.

```
# n is 10000, p is 100
bart.model <- bartMachine(X,y,
                           num_trees = 100,
                           num_burn_in = 1000,
                           num_iterations_after_burn_in = 5000,
                           mem_cache_for_speed = TRUE)
```

```
CPU [|||||] 100.0%   IO [|||||] 100.0%
```

```
Mem[|||||] 148G/189G
```

```
Swp[|||||] 0K/31.4G
```

- ▶ Statistical advantages are numerous
- ▶ Package authors aggregated many academic works on BART
- ▶ Great for small to mid sized data
- ▶ Good results with removing expected variation and feeding residuals into BART.